

2025 Introduction to Computer Software Systems Lab Assignment #2

Doyoung Kim

September 15, 2025

In this assignment, we implement fundamental operations in the two's complement integer and floating-point systems using the C programming language. In order to accomplish this, it requires the following six functions to be implemented under a specific set of constraints:

negate Implement the negation operator in the two's complement system.

isLess Compare two signed integers, returning true if the first argument is less than the second.

float_abs Compute the absolute value of a floating-point number using only integer operations.

float_twice Double a floating-point number by manipulating its bit-level representation.

float_i2f Implement the casting operation from an integer to a floating-point number.

float_f2i Implement the casting operation from a floating-point number to an integer.

We may assume a two's complement, 32-bit representation for the signed integer datatype and IEEE 754 compliance for single-precision floating-point numbers throughout this assignment. The bit-level representation of an IEEE 754 single-precision float is as follows:




Figure 1: Bit-level representation of an IEEE 754 single-precision floating-point number.

The numerical value of this representation is given by the formula:

$$(-1)^s \times M \times 2^{e-\text{bias}-p}$$

where s is the sign bit, M is the integral significand (mantissa), e is the exponent, and p is the precision of the integral significand. For normalized values, the significand M includes an implicit leading 1.

Special cases are determined by the exponent field. If the exponent is all ones, the value is either Not a Number (NaN) or infinity. If the exponent is all ones and the fraction is zero, the value represents infinity; otherwise, it is NaN. If the exponent field is all zeros, the value is denormalized and does not have an implicit leading 1.

This report will discuss the implementation details and the underlying principles of two's complement and floating-point arithmetic as they apply to modern computer systems.

1 Implementation

1.1 Function negate

This task requires the implementation of the two's complement unary negation operator. This is achieved by utilizing the principle of negation in a two's complement system, which states that the negation of an integer can be obtained by inverting all of its bits and adding one.

```
int negate(int x) { return ~x + 1; }
```

1.2 Function isLess

This function determines if the first argument is less than the second, returning 1 if true and 0 otherwise. A naive approach using the identity $x < y \iff x - y < 0$ is insufficient, as the subtraction may overflow in the two's complement system.

Therefore, this implementation first checks if the signs of the arguments differ. If they do, a negative number is always less than a positive one. If the signs are the same, the subtraction-based identity is safe to use, as overflow will not occur.

```
int isLess(int x, int y) {
    int sign_diff = ((x ^ y) >> 31) & 1;
    return (sign_diff & !(y >> 31)) |
           ((!sign_diff) & ((x + (~y + 1)) >> 31 & 1));
}
```

1.3 Function `float_abs`

The goal of this task is to compute the absolute value of a floating-point number by manipulating its bit-level representation. Since the most significant bit of a float represents its sign, this can be accomplished by clearing the sign bit.

A required edge case is that if the argument is NaN, the function must return the argument unmodified. Consequently, the implementation first checks for the NaN condition (exponent field is all ones and the fraction is non-zero) before proceeding.

```
unsigned float_abs(unsigned uf) {
    unsigned mask_exp = 0xFF << 23;

    /* Check if 'uf' is NaN. */
    if ((uf & mask_exp) == mask_exp && (uf & 0x007FFFFFFF) != 0)
        return uf;

    return 0x7FFFFFFF & uf;
}
```

1.4 Function `float_twice`

This task requires doubling a floating-point number. The primary strategy is to add one to the exponent field. However, several special cases require careful consideration.

First, if the input is NaN or infinity, it should be returned unchanged. Second, if the input is a denormalized number (including zero), adding one to the exponent is incorrect. Instead, the entire number (excluding the sign bit) is shifted left by one. Third, if incrementing the exponent of a normalized number results in an exponent of all ones, the number has overflowed to infinity, and its fraction must be cleared to prevent it from being misinterpreted as a NaN.

```
unsigned float_twice(unsigned uf) {
    unsigned mask_exp = 0xFF << 23;
    unsigned exp = uf & mask_exp;

    if (exp == mask_exp)
        return uf;

    if (exp == 0)
        return (uf & 0x80000000) | (uf << 1);

    exp += 1 << 23;

    if (exp == mask_exp)
        return (uf & (1 << 31)) | mask_exp;
    else
```

```

        return (uf & ~mask_exp) | exp;
    }

```

1.5 Function float_i2f

Converting an integer to a floating-point number is a multi-step process that involves determining the final sign, exponent, and fraction, as well as performing rounding for values that cannot be represented precisely.

This implementation correctly determines the components and then performs rounding according to the *round-to-nearest, ties-to-even* principle. It examines the round bit (the first bit shifted out) and the sticky bit (the logical OR of all subsequent shifted-out bits) to decide if the fraction needs to be incremented. In a tie situation, the final bit of the fraction (the guard bit) is checked to ensure the result is even.

```

unsigned float_i2f(int x) {
    unsigned frac, exp, round, sticky, increment = 0;

    if (x == 0)
        return 0;

    if (x < 0)
        frac = -x;
    else
        frac = x;

    exp = 159;

    while ((frac & 0x80000000) == 0) {
        frac <<= 1;
        exp--;
    }

    frac <<= 1;
    exp--;

    round = frac & 0x100;
    sticky = frac & 0xFF;

    if (round && !sticky)
        increment = (frac & 0x200) != 0;
    else if (round && sticky)
        increment = 1;

    frac >>= 9;
    return ((x & 0x80000000) | (exp << 23) | frac) + increment;
}

```

```
}
```

1.6 Function float_f2i

The conversion from a floating-point number to an integer involves reconstructing the integer value from the float's sign, exponent, and fraction. The value can be reconstructed by shifting the implicit '1' and the fractional bits left by a number of positions dictated by the exponent.

A key aspect of this conversion in C is that it performs *truncation* (rounding toward zero). Therefore, no rounding logic is required. After the integer part is reconstructed, the final result is obtained by negating it if the original float was negative. The implementation also correctly handles overflow by returning '0x80000000' for values outside the representable range of a 32-bit signed integer.

```
int float_f2i(unsigned uf) {
    int exp;
    unsigned integral, fractional;

    if ((uf & 0x7FFFFFFF) == 0)
        return 0;

    exp = ((uf & 0x7F800000) >> 23) - 127;

    if (exp < 0)
        return 0;

    if (exp >= 31)
        return 0x80000000;

    integral = 1 << exp;
    integral |= (uf & 0x007FFFFF) >> (23 - exp);

    if (uf & 0x80000000)
        return -integral;
    else
        return integral;
}
```

2 Result and Discussion

The execution of the `driver.pl` script confirms that all implemented functions pass the required correctness and performance tests, achieving a perfect score.

```
[d0319@programming2 datalab2]$ ./driver.pl
...
```

Score = 31/31 [19/19 Corr + 12/12 Perf] (80 total operators)

The results validate the correctness of the bit-level manipulation strategies employed for both two's complement integers and IEEE 754 floating-point numbers. The most complex functions, `float_i2f` and `float_f2i`, successfully implemented intricate behaviors like round-to-nearest-even and truncation, respectively, demonstrating a comprehensive understanding of data representation.

3 Conclusion

This assignment was successfully completed, with all functions for manipulating integer and floating-point representations meeting the specified correctness and performance criteria. The lab provided significant practical insight into the low-level mechanics of computer arithmetic.

The primary takeaway from this exercise is the critical role of bit-level representation in defining the behavior of numerical operations. Successfully implementing functions like `float_twice` and `float_i2f` required moving beyond abstract numerical values to directly manipulating their underlying sign, exponent, and fraction fields. This process reinforces the idea that for a computer, numbers are not just mathematical concepts but structured patterns of bits governed by precise rules.

Furthermore, the assignment highlighted the nuances and edge cases inherent in computer arithmetic, such as handling NaNs, infinities, denormalized numbers, and overflow conditions. Mastering these details is fundamental to writing robust, reliable systems-level code. Ultimately, this lab serves as an effective bridge between the theoretical principles of data representation and their practical application in software.